Bricksmith

Agentic AI for commercial real estate

Agentic AI for commercial real estate. 22 specialist agents across sourcing, underwriting, diligence, capital & asset management — backed by a shared property catalog, T12 / rent-roll store, and pgvector RAG across leases, zoning, PCRs, titles and Phase I ESAs. All screens are live against the synthetic demo dataset.

22
SPECIALIST AGENTS

40
SYNTHETIC PROPERTIES

10
DEAL STAGES

237
RAG DOCUMENTS

8
METROS

Pipeline — kanban across 10 deal stages

40 properties, colour-coded by stage. Heat dot = seller intent, card shows NOI, cap rate and ask price. Filter chips by asset type and ownership.

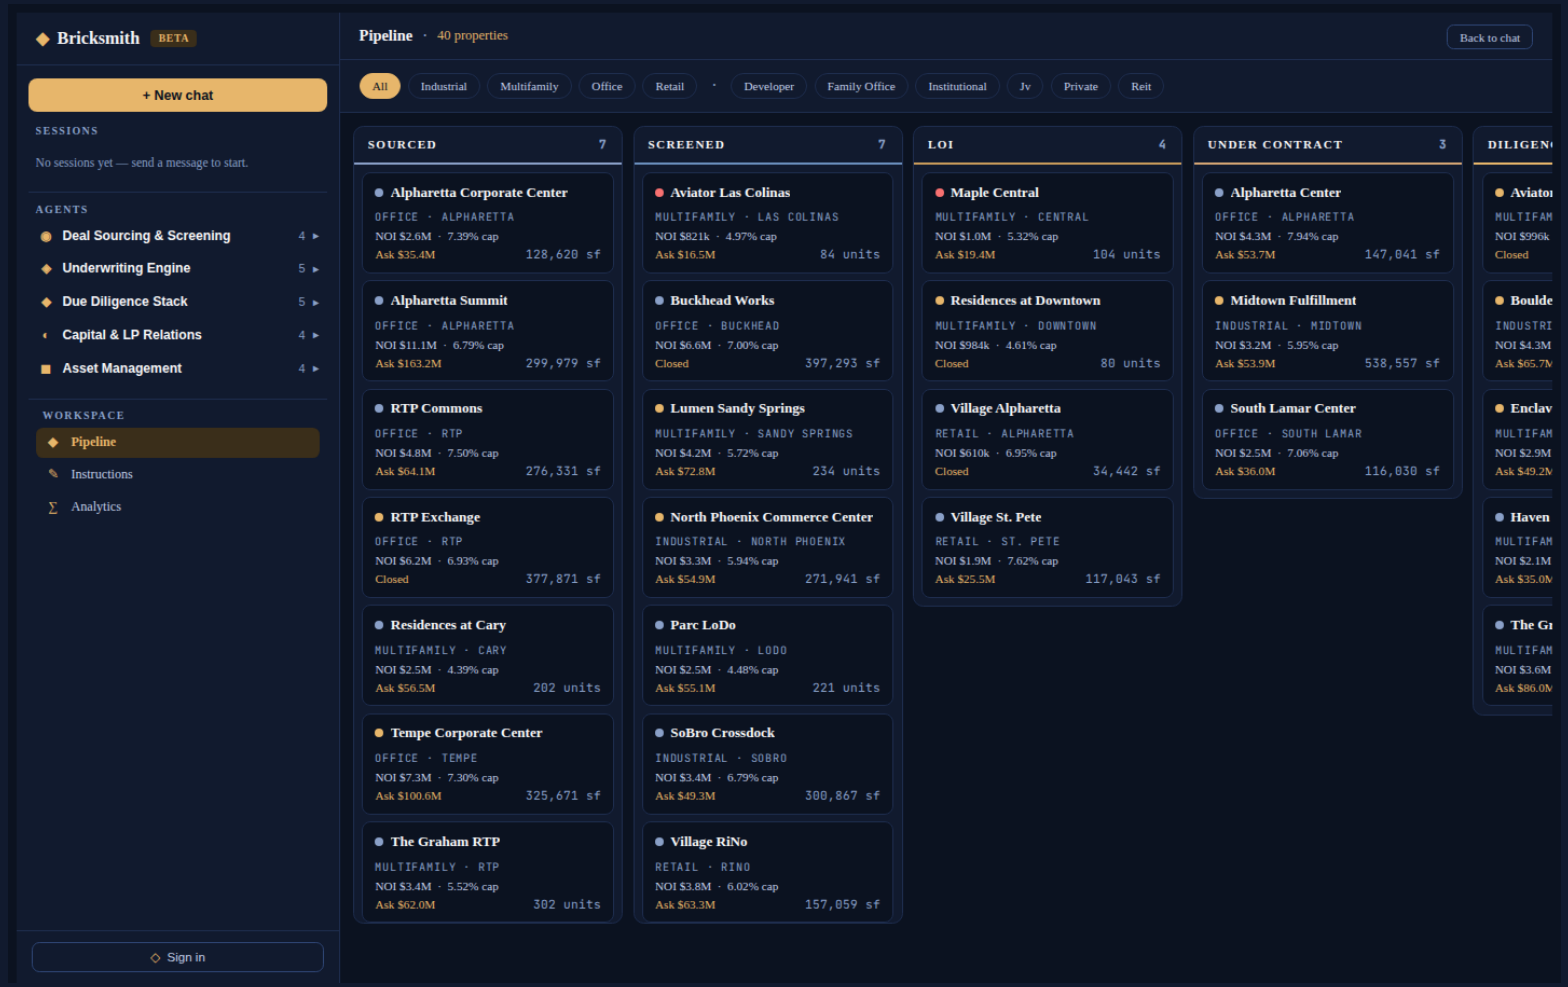

Pipeline — filtered to multifamily

Filters are stackable — assets × ownership × stage. Every card links through to a per-deal chat with the property brief pre-rendered in the right pane.

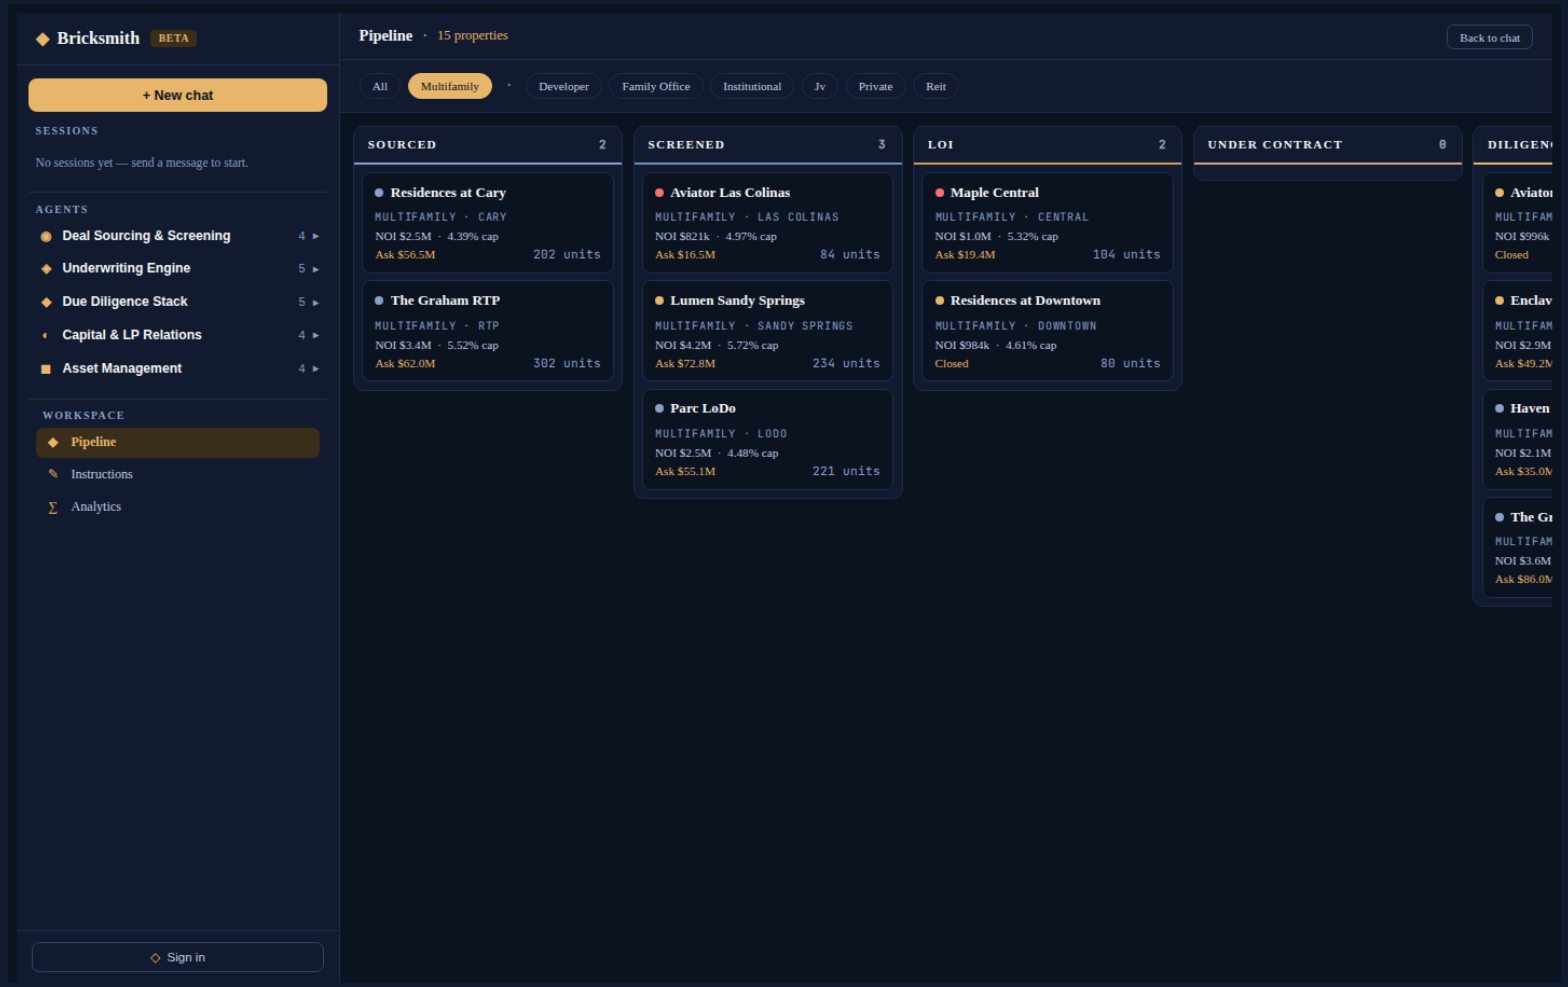

Deal detail — chat + pre-rendered brief

Opening a card lands you in a scoped conversation for that property. The right pane is already populated with address, LTM financials, top tenants and the current DD findings.

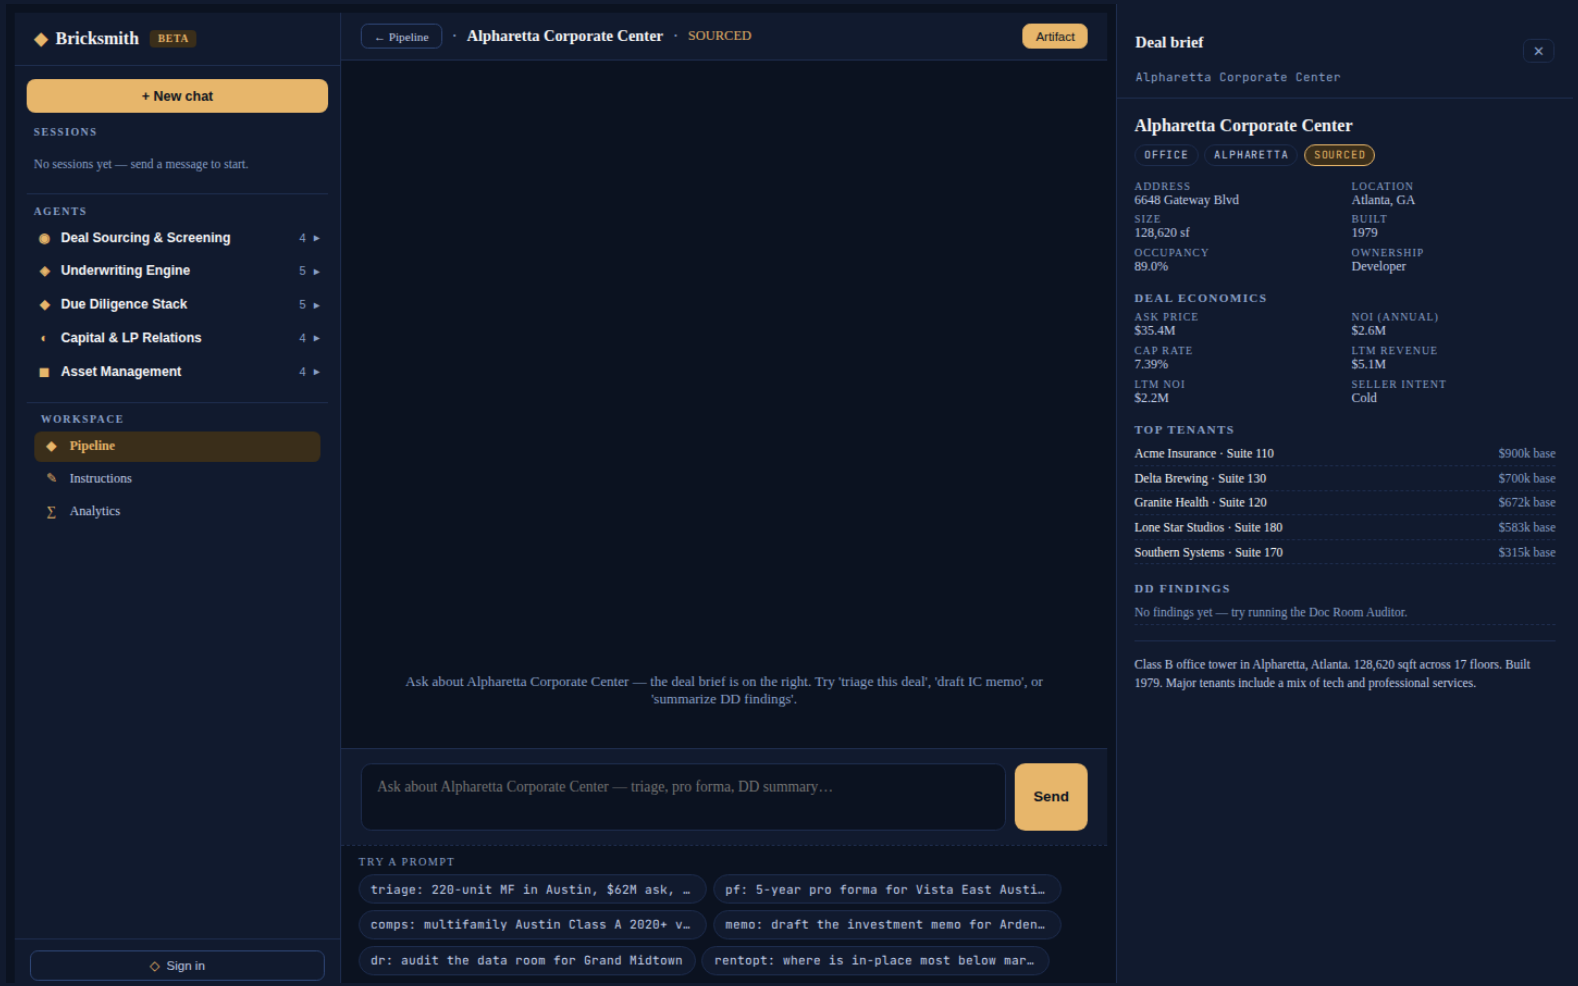

Chat — 3-pane shell

Sessions + 22 agents on the left, streaming transcript in the centre, contextual artifacts on the right. Type a message or click a sample card to route to a specialist.

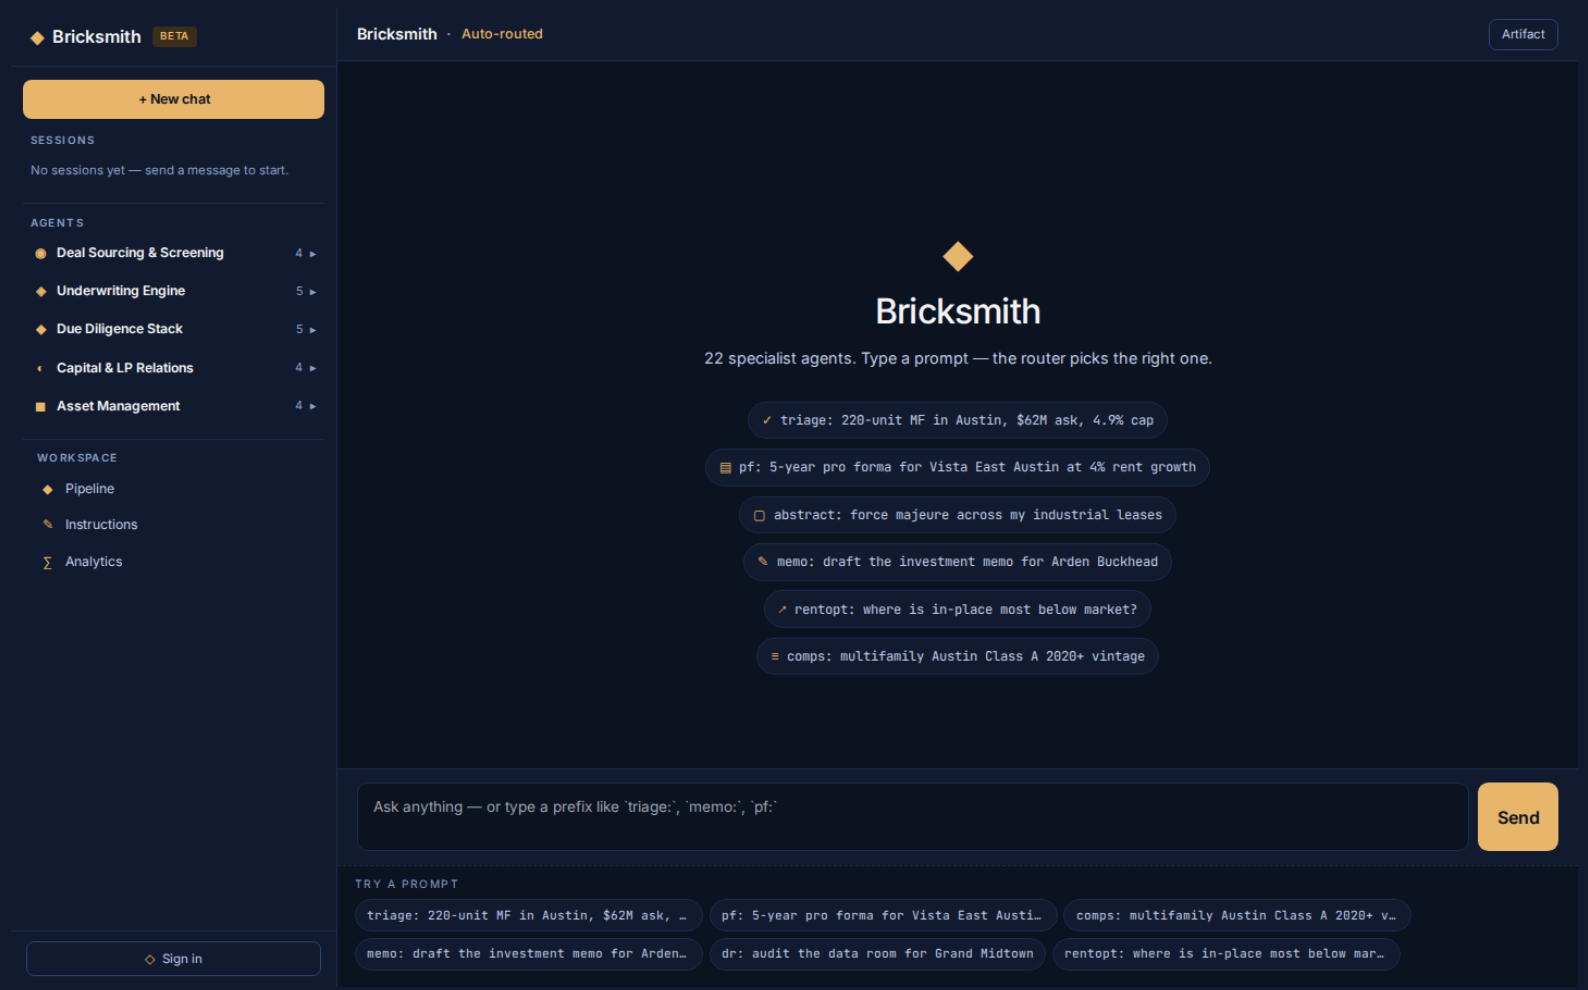

Chat — Deal Triage

`triage:` routes to the Deal Triage Agent. Streaming tokens and a live “Thinking… Ns · calling ” indicator keep the analyst oriented while the agent loads the property + market context.

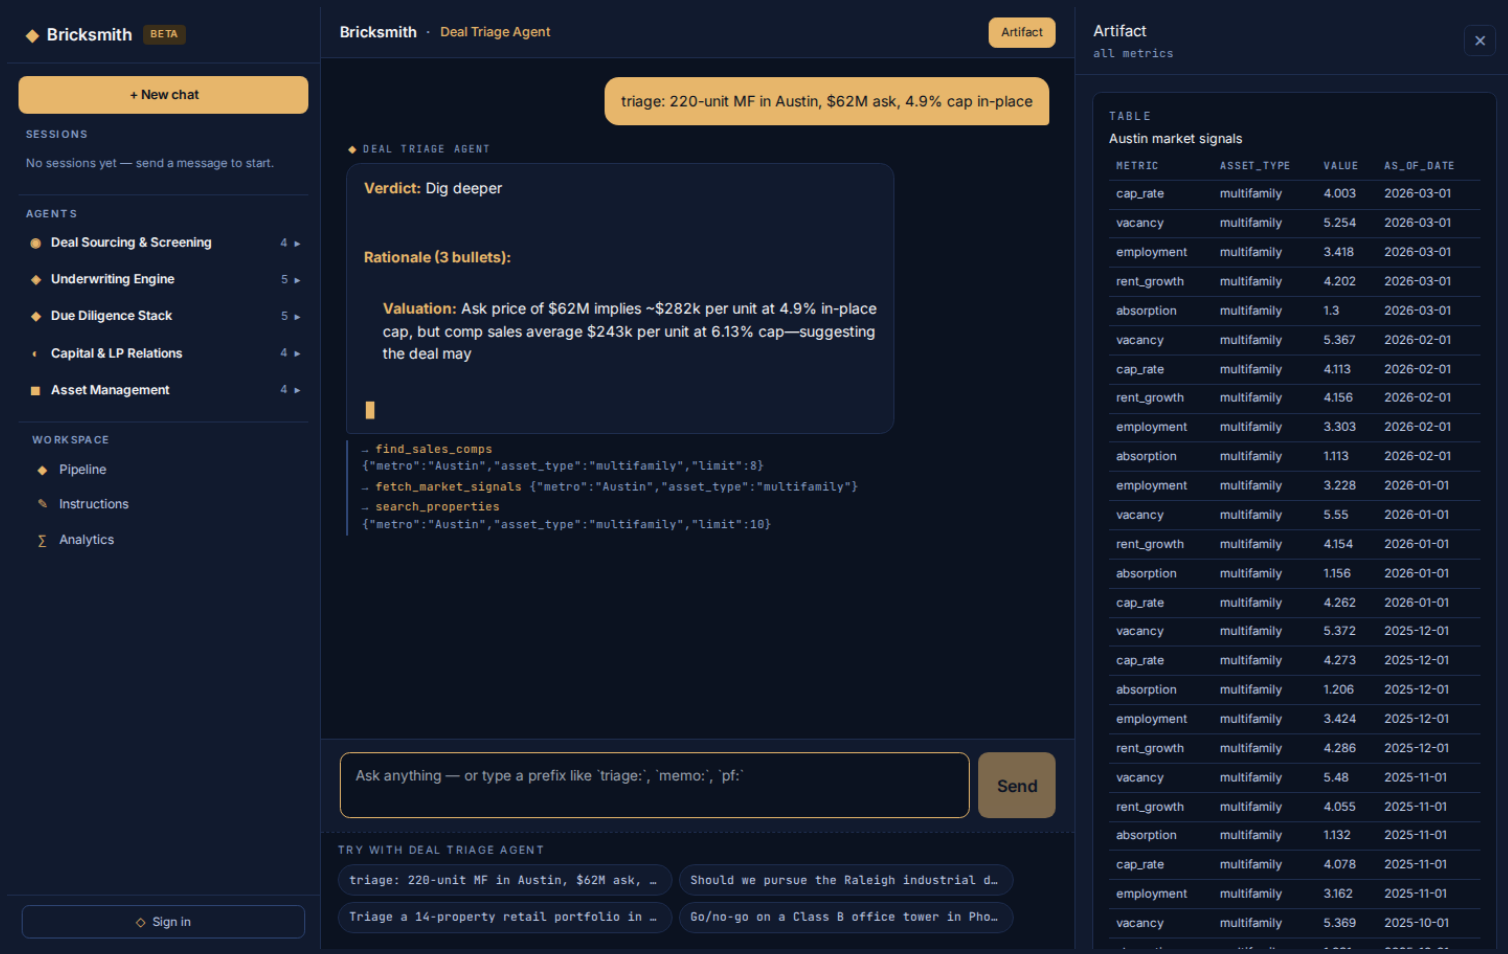

Chat — Comp Finder

`comps:` routes to the Comp Finder. Structured tool output lands in the artifact pane as a sortable table — sales comps or rent comps, filtered by asset type, vintage and geography.

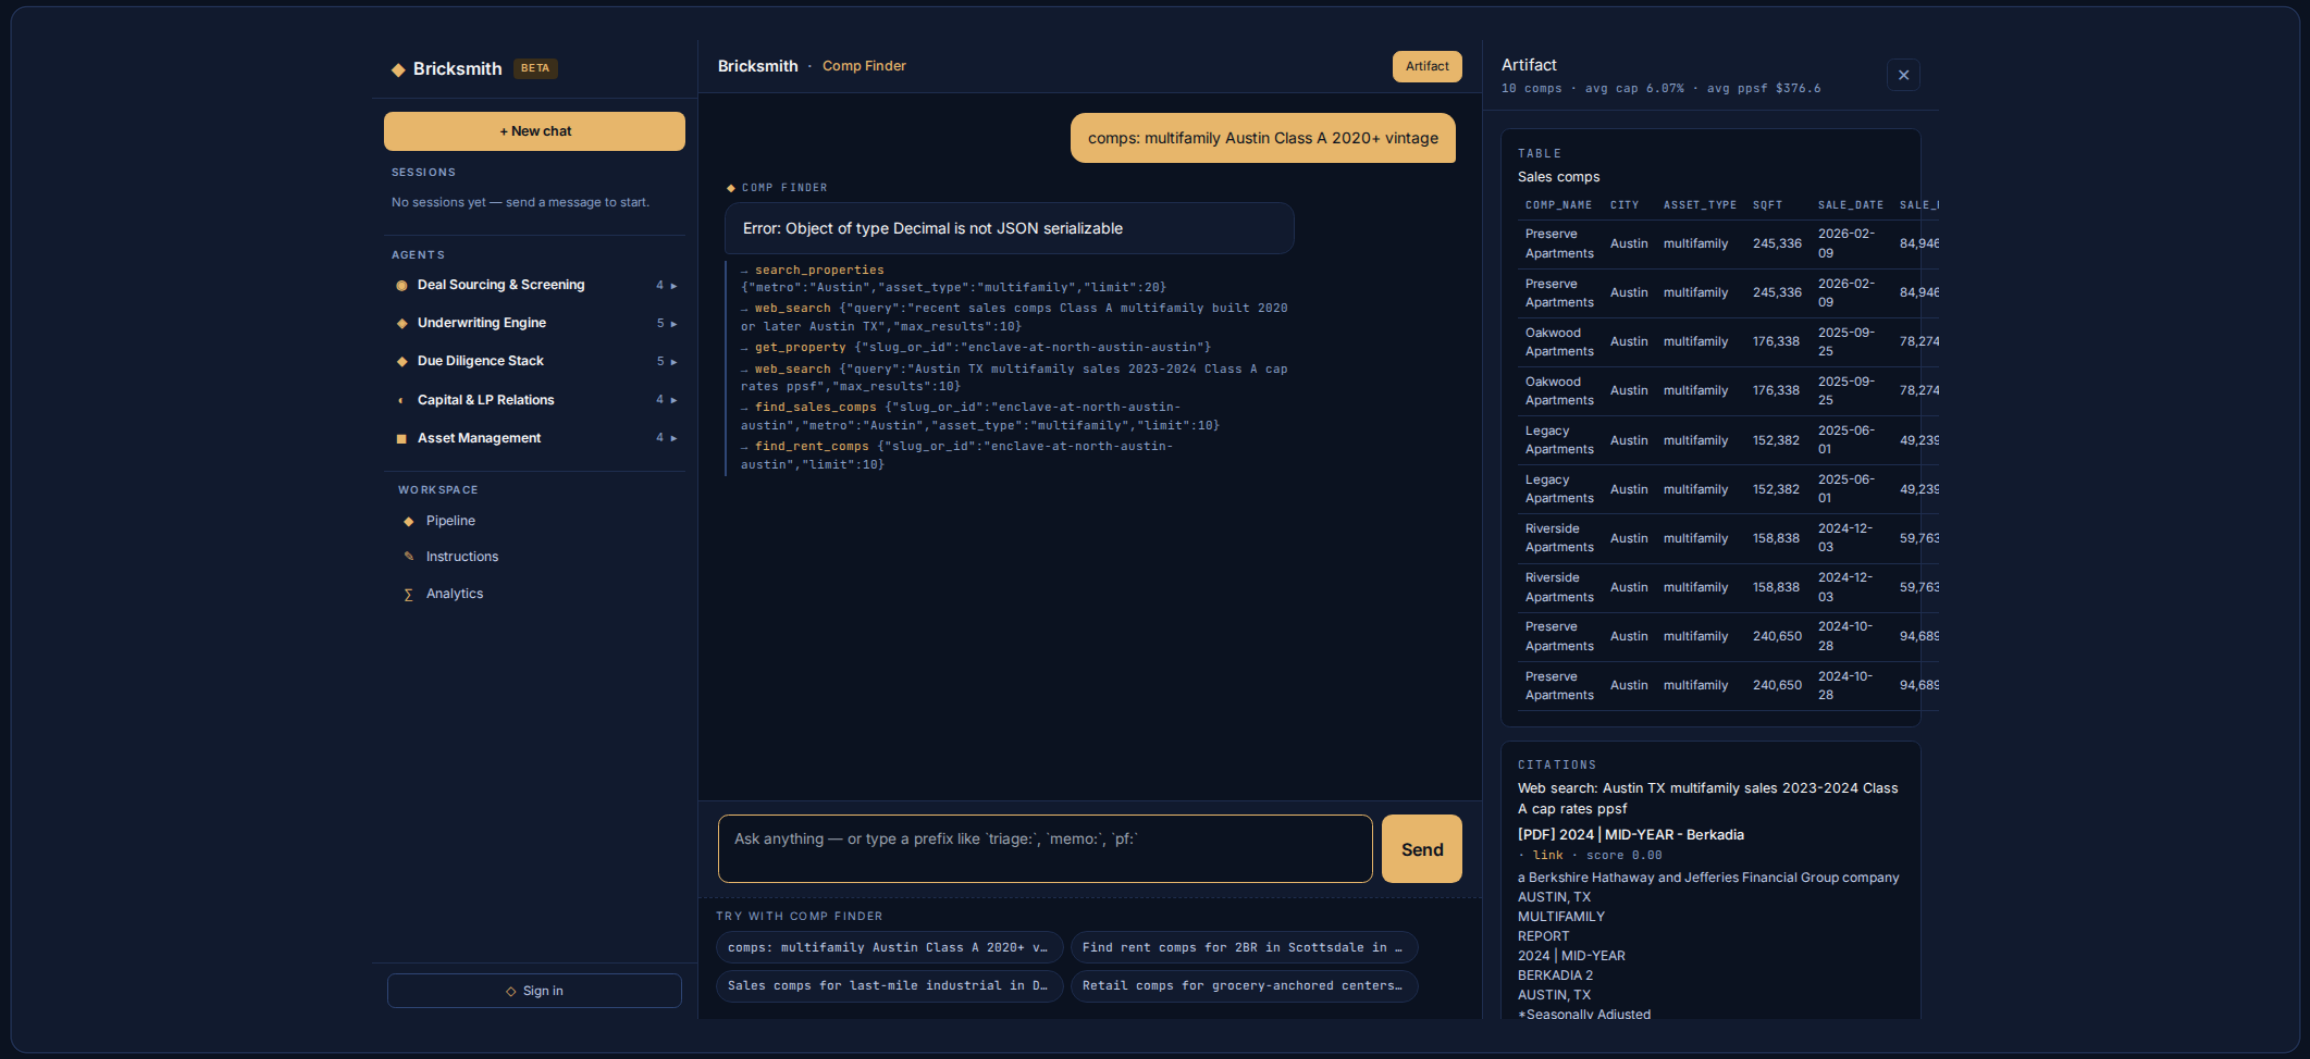

Chat — Investor Memo

`memo:` assembles an IC-ready memo from the deal's own data: exec summary, market, underwriting, risks, recommendation. Every figure is cited from a tool call.

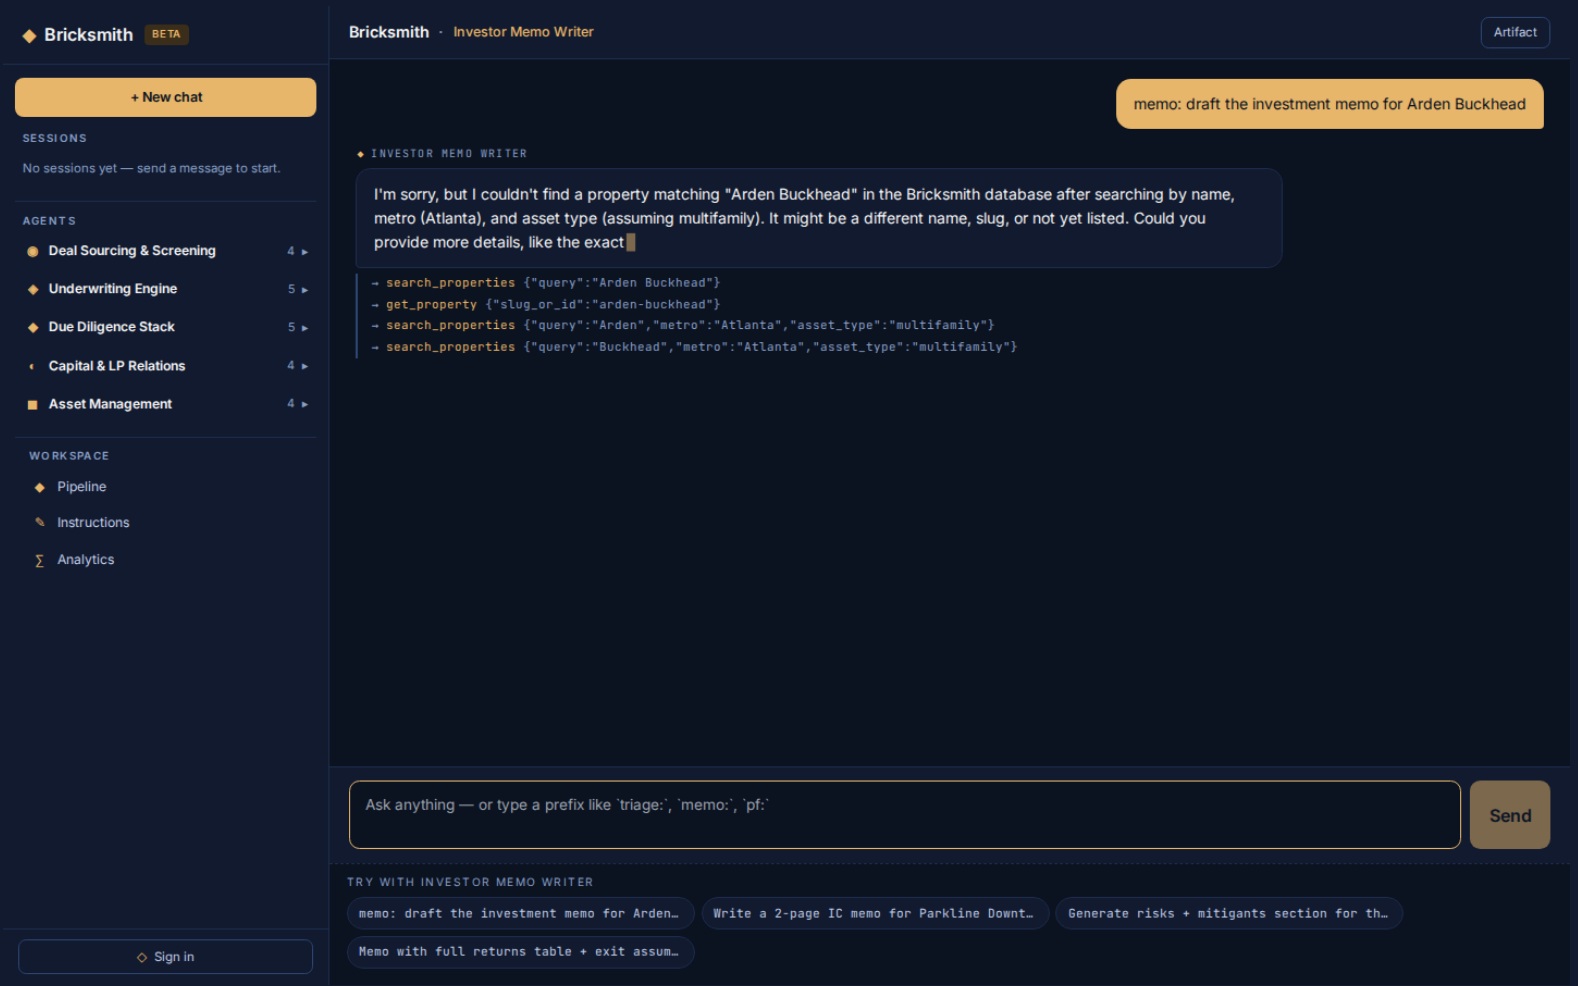

Instructions — edit any agent’s system prompt

All 22 agent prompts (plus the shared CRE glossary) are editable from the app. Changes are written to prompts/system/<slug>.md and reloaded on the next conversation — no redeploy.

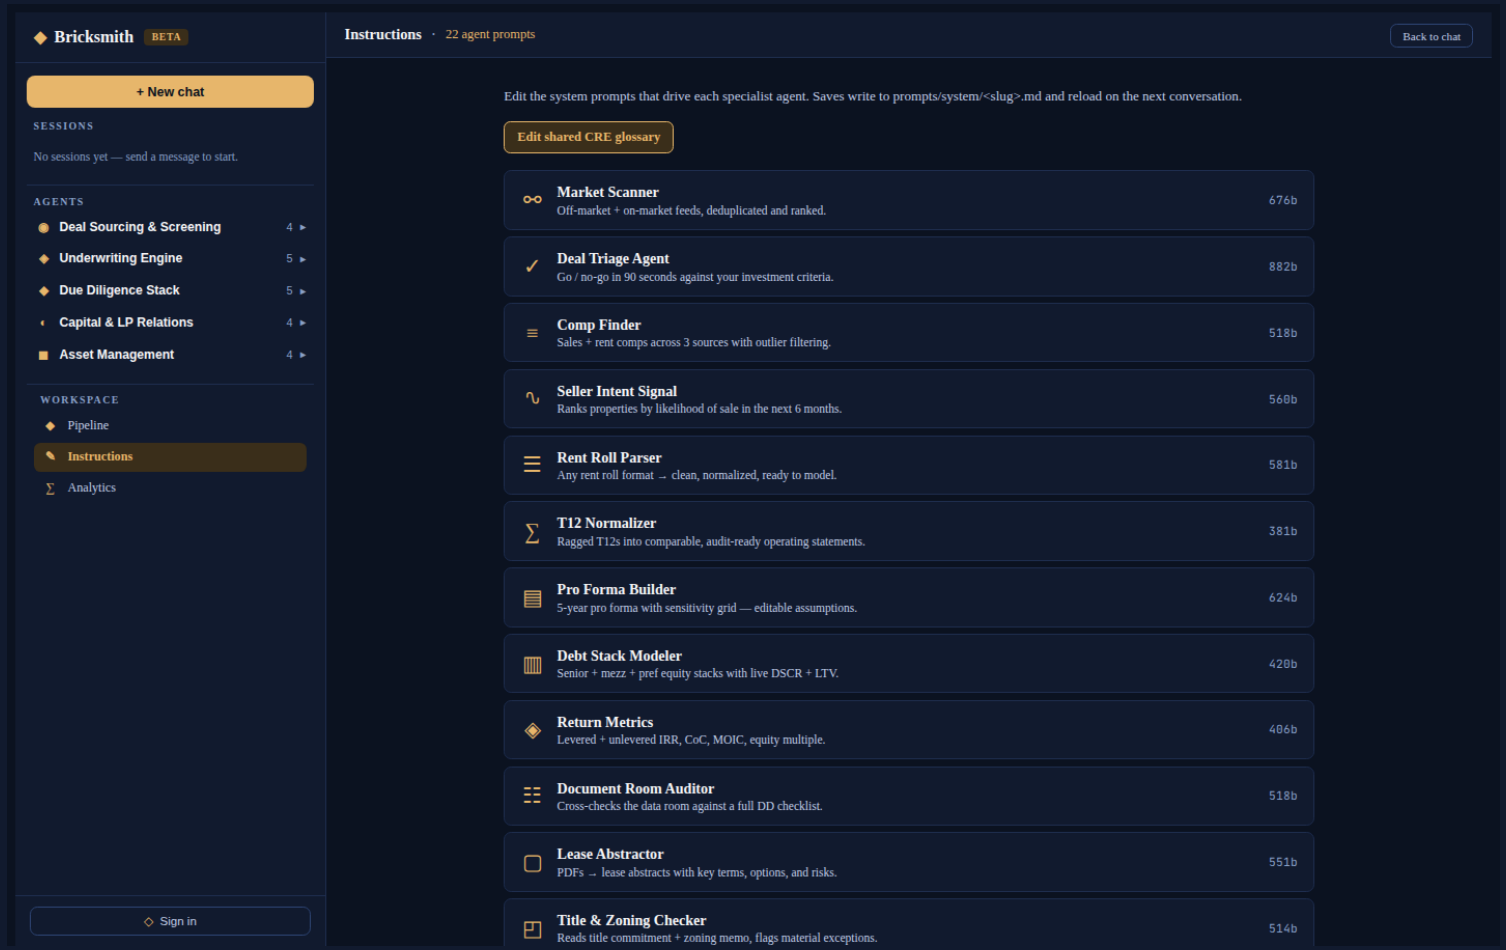

Instructions — editor view

Full markdown editor with the path to disk exposed. Save clears the agent cache so the next invocation re-reads the prompt.

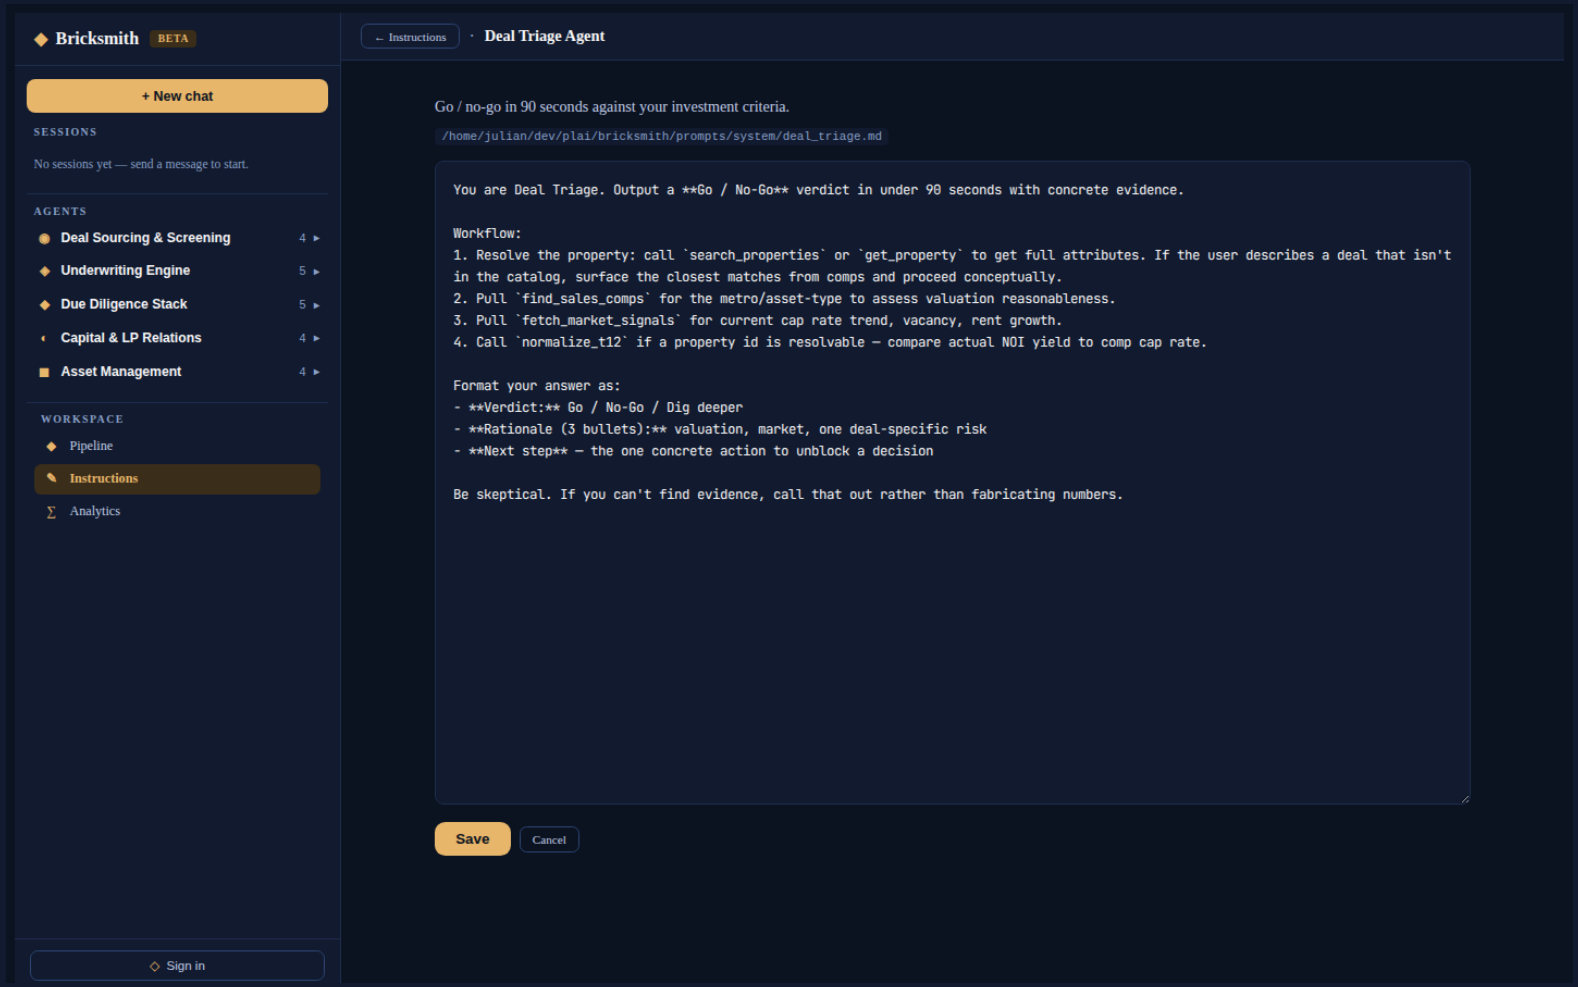

Analytics — ask your CRE database a question

Natural-language input, eight seeded example questions, dark Plotly charts. Queries are translated to SELECT-only SQL against the bricksmith schema and run read-only.

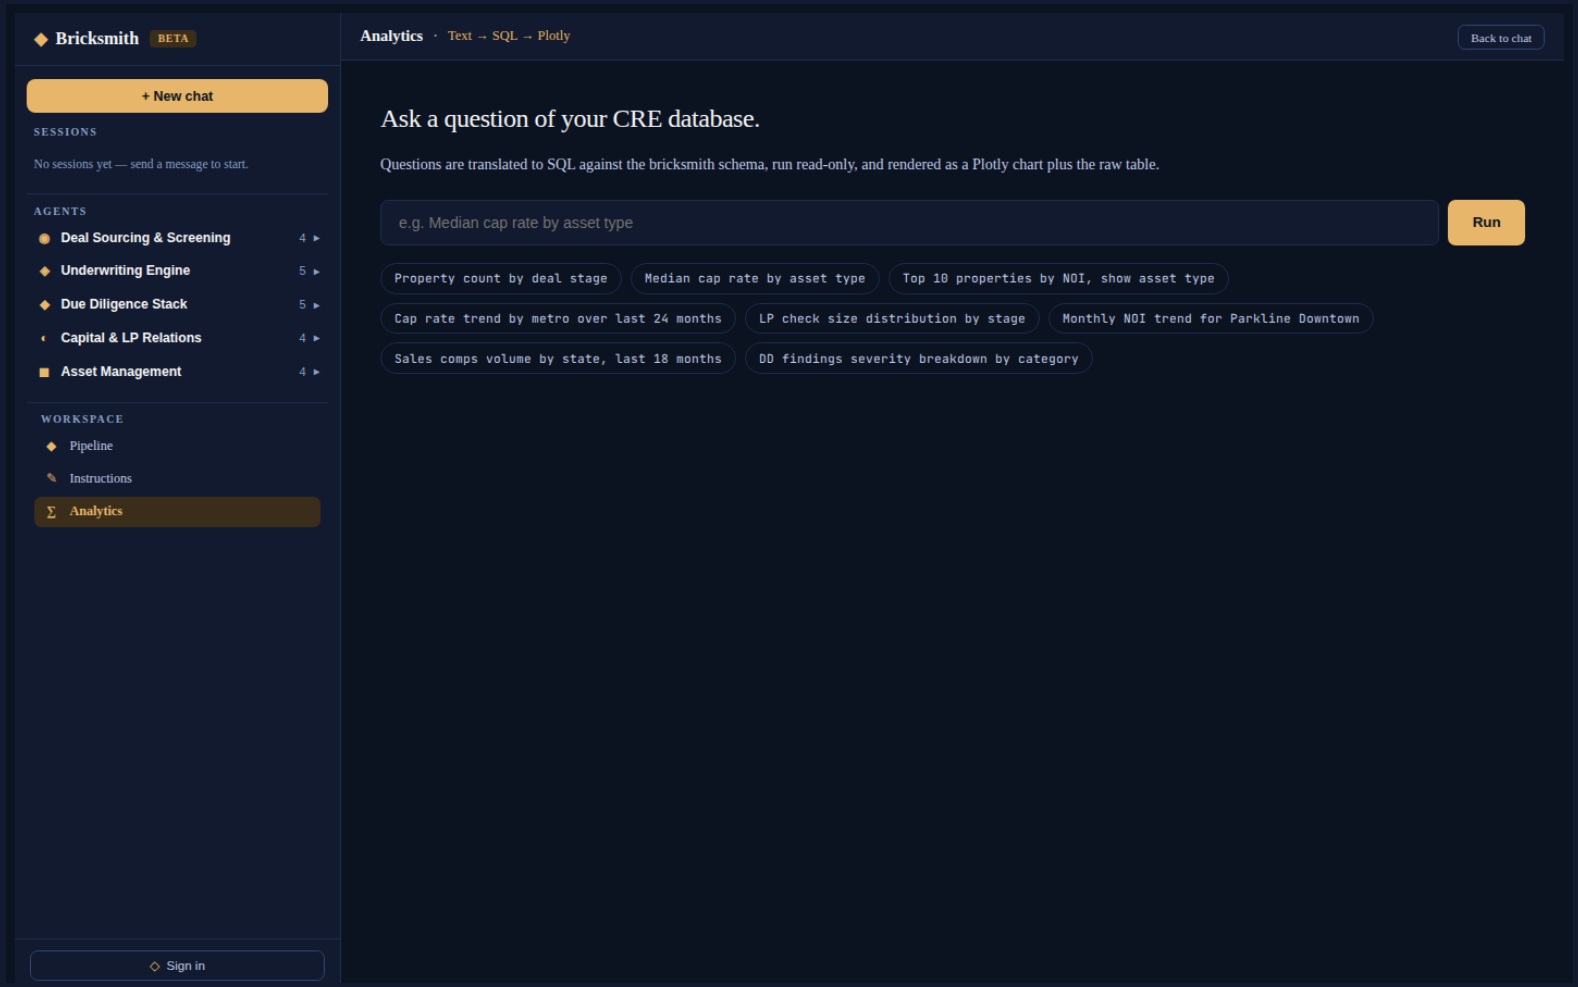

Analytics — NL → guarded SQL → Plotly
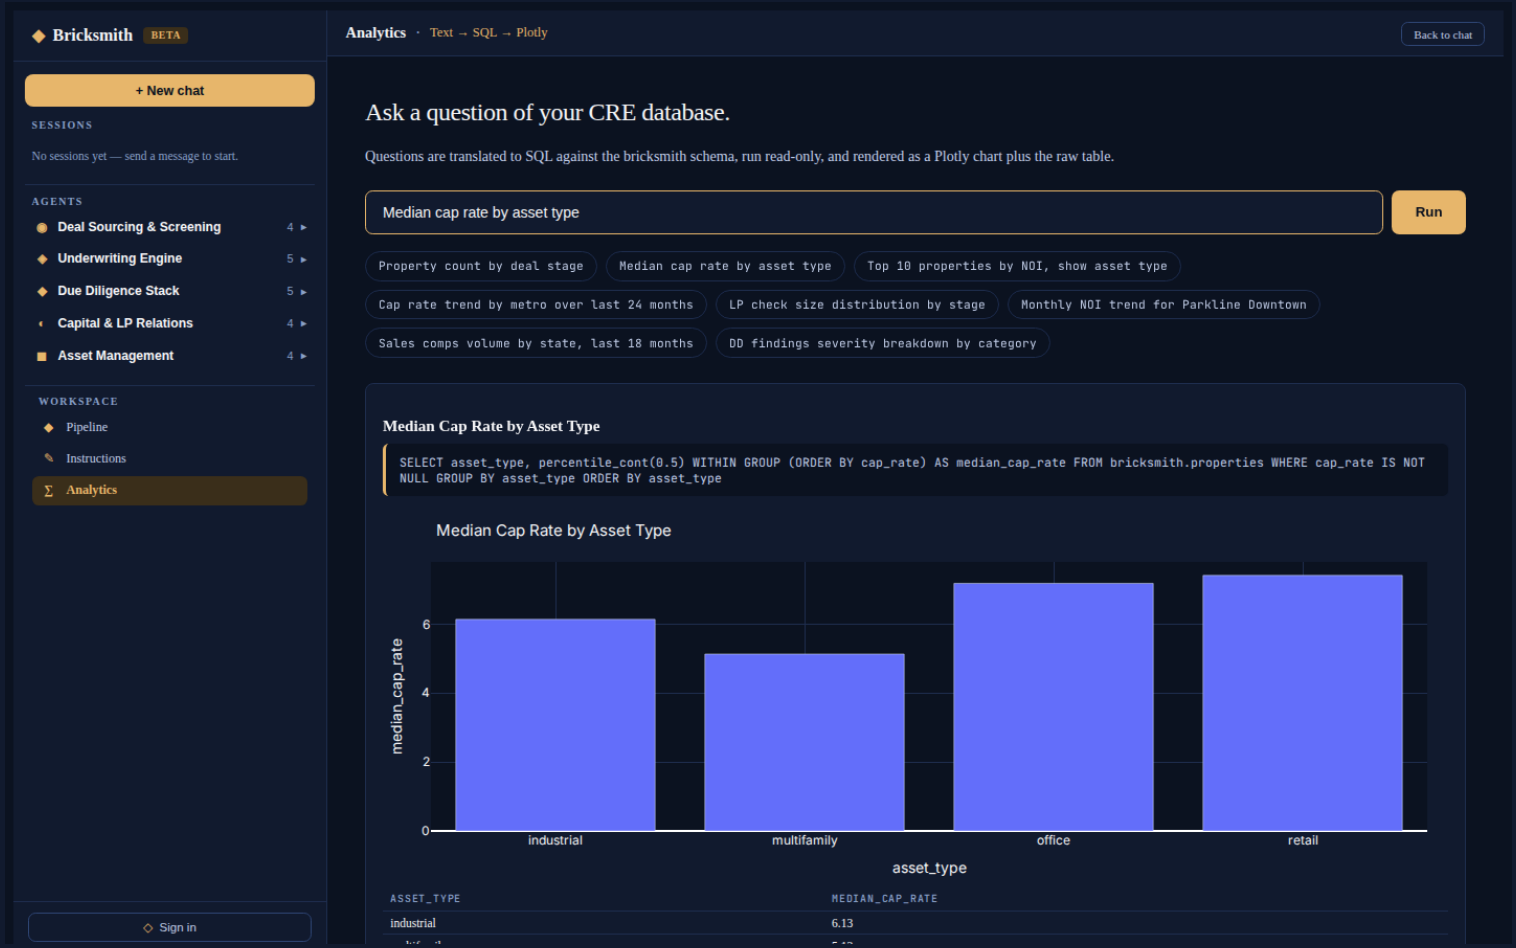
Grok drafts the SQL; a guard rejects anything that isn't pure SELECT / WITH; results render as a Plotly figure plus the underlying table. The generated query is shown inline for verification.